

✓ Kaggle GAN Competition: Dog Images

✓ Challenge Problem:

In this task, we aim to generate realistic data samples from a complex and high-dimensional true data distribution, which is often inaccessible. To achieve this, we will employ a Generative Adversarial Network (GAN), consisting of two neural networks: a Generator G and a Discriminator D. The Generator creates synthetic data samples by transforming random noise into data resembling the true distribution, while the Discriminator distinguishes between real and fake samples. These networks are trained simultaneously in a game-theoretic manner, where the Generator aims to fool the Discriminator, and the Discriminator works to correctly identify real versus fake data. This adversarial process improves the performance of both networks, gradually allowing the Generator to produce more realistic data.

Source # <https://www.kaggle.com/competitions/gan-getting-started/overview>

```
import warnings
import os
import time
from glob import glob
import datetime
import seaborn as sns
import numpy as np, pandas as pd, os
import keras
import matplotlib.pyplot as plt, zipfile
from PIL import Image

import zipfile
from keras.layers import Input, Dense, Reshape, Flatten, Dropout, Concatenate
from tensorflow.keras.backend import random_normal, ones_like, zeros_like, mean
from keras.layers import Conv2D, UpSampling2D
from keras.layers import BatchNormalization, Activation, ZeroPadding2D
from keras.models import Sequential, Model
from keras.optimizers import Adam
from keras.layers import concatenate
from keras.initializers import TruncatedNormal
from keras.callbacks import LearningRateScheduler, EarlyStopping, History
from PIL import Image
import warnings
import xml.etree.ElementTree as ET
import os
print(os.listdir("/kaggle/input/generative-dog-images"))
```

```
↳ ['Annotation.zip', 'all-dogs.zip']
```

```
with zipfile.ZipFile("../input/generative-dog-images/all-dogs.zip", "r") as z:
    z.extractall("/kaggle/working/")

dog_dir = "/kaggle/working/all-dogs/"
filenames = os.listdir(dog_dir)

print('Total images:', len(filenames))

with zipfile.ZipFile("../input/generative-dog-images/Annotation.zip", "r") as z:
    z.extractall("/kaggle/working/")

dog_dir_Annotation = "/kaggle/working/Annotation/"
filenames_Annotation = os.listdir(dog_dir_Annotation)

print('Total images:', len(filenames_Annotation))

↳ Total images: 20579
   Total images: 120

ComputeLB = False
DogsOnly = False

ROOT = '/kaggle/working/'
if not ComputeLB: ROOT = '/kaggle/working/'
```

```

IMAGES = os.listdir(ROOT + 'all-dogs/')
breeds = os.listdir(ROOT + 'Annotation/')

idxIn = 0; namesIn = []
imagesIn = np.zeros((25000,64,64,3))

# CROP WITH BOUNDING BOXES TO GET DOGS ONLY
# https://www.kaggle.com/paulorzp/show-annotations-and-breeds
if DogsOnly:
    for breed in breeds:
        for dog in os.listdir(ROOT+'Annotation/'+breed):
            try: img = Image.open(ROOT+'all-dogs/'+dog+'.jpg')
            except: continue
            tree = ET.parse(ROOT+'Annotation/'+breed+'/'+dog)
            root = tree.getroot()
            objects = root.findall('object')
            for o in objects:
                bndbox = o.find('bndbox')
                xmin = int(bndbox.find('xmin').text)
                ymin = int(bndbox.find('ymin').text)
                xmax = int(bndbox.find('xmax').text)
                ymax = int(bndbox.find('ymax').text)
                w = np.min((xmax - xmin, ymax - ymin))
                img2 = img.crop((xmin, ymin, xmin+w, ymin+w))
                img2 = img2.resize((64,64), Image.ANTIALIAS)
                imagesIn[idxIn,:,:,:] = np.asarray(img2)
                #if idxIn%1000==0: print(idxIn)
                namesIn.append(breed)
                idxIn += 1
    idx = np.arange(idxIn)
    np.random.shuffle(idx)
    imagesIn = imagesIn[idx,:,:,:]
    namesIn = np.array(namesIn)[idx]


# RANDOMLY CROP FULL IMAGES
else:
    IMAGES = np.sort(IMAGES)
    np.random.seed(810)
    x = np.random.choice(np.arange(20579),10000)
    np.random.seed(None)
    for k in range(len(x)):
        img = Image.open(ROOT + 'all-dogs/' + IMAGES[x[k]])
        w = img.size[0]; h = img.size[1];
        if (k%2==0)|(k%3==0):
            w2 = 100; h2 = int(h/(w/100))
            a = 18; b = 0
        else:
            a=0; b=0
            if w<h:
                w2 = 64; h2 = int((64/w)*h)
                b = (h2-64)//2
            else:
                h2 = 64; w2 = int((64/h)*w)
                a = (w2-64)//2
        img = img.resize((w2,h2), Image.Resampling.LANCZOS)
        img = img.crop((0+a, 0+b, 64+a, 64+b))
        imagesIn[idxIn,:,:,:] = np.asarray(img)
        namesIn.append(IMAGES[x[k]])
        #if idxIn%1000==0: print(idxIn)
        idxIn += 1

# DISPLAY CROPPED IMAGES
x = np.random.randint(0,idxIn,25)
for k in range(5):
    plt.figure(figsize=(15,3))
    for j in range(5):
        plt.subplot(1,5,j+1)
        img = Image.fromarray( imagesIn[x[k*5+j],,:,:].astype('uint8') )
        plt.axis('off')
        if not DogsOnly: plt.title(namesIn[x[k*5+j]],fontsize=11)
        else: plt.title(namesIn[x[k*5+j]].split('-')[1],fontsize=11)
        plt.imshow(img)
    plt.show()

```




n02109525_5385.jpg



n02090721_4073.jpg




n02093754_3749.jpg



n02088094_2458.jpg



n02091635_4174.jpg




n02113978_3640.jpg




n02093754_897.jpg



n02097474_4929.jpg




n02093256_1371.jpg



n02113186_10416.jpg




n02091244_2038.jpg




n02107312_3841.jpg



n02093428_3397.jpg



n02105251_7977.jpg




n02105855_5165.jpg




n02112137_8983.jpg




n02101556_4386.jpg



n02101006_3354.jpg




n02086646_1342.jpg




n02111129_4069.jpg




n02091032_12013.jpg




n02086646_166.jpg




n02090721_2347.jpg



n02097130_1213.jpg



n02099267_1272.jpg



```
IMG_SIZE = Input((12288,))
IMG_SIZE_2 = Input((10000,))
NOISE_SIZE = 10000
#BATCH_SIZE = 256 # orig gives ~7.24
#BATCH_SIZE = 512 # gives ~7.25
BATCH_SIZE = 128 # gives 7.22594
#BATCH_SIZE = 128 # gives ~7.222
#BATCH_SIZE = 64
```

✓ Build Discriminator

```
def discriminatorFunction():  
    input_layer = Dense(12288, activation='sigmoid')(IMG_SIZE_2)  
    input_layer = Reshape((2,12288,1))(concatenate([IMG_SIZE,input_layer]))  
    discriminator = Conv2D(filters = 1, kernel_size=[2,1],use_bias=False, name = 'layer_1')(input_layer)  
    out = Flatten()(discriminator)  
    return out  
  
print("Discriminator")  
model = discriminatorFunction()  
model_discriminator = Model([IMG_SIZE,IMG_SIZE_2], model)  
model_discriminator.get_layer('layer_1').trainable = False  
model_discriminator.get_layer('layer_1').set_weights([np.array([[[[-1.0 ]],[[1.0]]]])])  
model_discriminator.summary()  
model_discriminator.compile(optimizer='adam', loss='binary_crossentropy')
```

🔗 Discriminator
I0000 00:00:1745788134.278856 31 gpu_device.cc:2022] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 13942 MB memc
I0000 00:00:1745788134.279680 31 gpu_device.cc:2022] Created device /job:localhost/replica:0/task:0/device:GPU:1 with 13942 MB memc
Model: "functional"

Layer (type)	Output Shape	Param #	Connected to
input_layer_3 (InputLayer)	(None, 10000)	0	-
input_layer_2 (InputLayer)	(None, 12288)	0	-
dense (Dense)	(None, 12288)	122,892,288	input_layer_3[0][0]
concatenate (Concatenate)	(None, 24576)	0	input_layer_2[0][0], dense[0][0]
reshape (Reshape)	(None, 2, 12288, 1)	0	concatenate[0][0]
layer_1 (Conv2D)	(None, 1, 12288, 1)	2	reshape[0][0]
flatten (Flatten)	(None, 12288)	0	layer_1[0][0]

Total params: 122,892,290 (468.80 MB)
Trainable params: 122,892,288 (468.80 MB)

✓ Build Generator

```
def GeneratorFunction(noise_shape=(NOISE_SIZE,)):  
    input_layer = Input(noise_shape)  
    generated = Dense(12288, activation='linear')(input_layer)  
  
# COMPILE  
    model = Model(inputs=input_layer,outputs = [generated,Reshape((10000,))(input_layer)])  
    model.summary()  
  
    return model  
  
print("Generator")  
model_generator = GeneratorFunction(noise_shape=(NOISE_SIZE,))
```

Generator
Model: "functional_1"

Layer (type)	Output Shape	Param #	Connected to
input_layer_6 (InputLayer)	(None, 10000)	0	-
dense_1 (Dense)	(None, 12288)	122,892,288	input_layer_6[0][0]
reshape_1 (Reshape)	(None, 10000)	0	input_layer_6[0][0]

Total params: 122,892,288 (468.80 MB)
Trainable params: 122.892.288 (468.80 MB)

✓ Training the Discriminator

```
# TRAINING DATA
train_y = (imagesIn[:10000,:,:,:]/255.).reshape((-1,12288))
train_X = np.zeros((10000,10000))
for i in range(10000): train_X[i,i] = 1
zeros = np.zeros((10000,12288))

# -----
# Train Discriminator
# -----

lr = 0.5
for k in range(5):
    LR_Scheduler = LearningRateScheduler(lambda x: lr)
    h = model_discriminator.fit([zeros,train_X], train_y, epochs = 100,batch_size = BATCH_SIZE, callbacks=[LR_Scheduler], verbose=0)
    print('Epoch',(k+1)*10,'/50 - loss =',h.history['loss'][-1] )
    if h.history['loss'][-1]<0.533: lr = 0.1
```

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
I0000 00:00:1745788173.749432 98 service.cc:148] XLA service 0x7f85d80090b0 initialized for platform CUDA (this does not guarantee
I0000 00:00:1745788173.751311 98 service.cc:156] StreamExecutor device (0): Tesla T4, Compute Capability 7.5
I0000 00:00:1745788173.751346 98 service.cc:156] StreamExecutor device (1): Tesla T4, Compute Capability 7.5
I0000 00:00:1745788174.044371 98 cuda_dnn.cc:529] Loaded cuDNN version 90300
I0000 00:00:1745788176.815428 98 device_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifet
Epoch 10 /50 - loss = 0.5325140357017517
Epoch 20 /50 - loss = 0.5324963927268982
Epoch 30 /50 - loss = 0.5324907302856445
Epoch 40 /50 - loss = 0.5324869155883789
Epoch 50 /50 - loss = 0.5324838757514954

```
del train_X, train_y, imagesIn
```

✓ Building GAN

We will purposely give our Generator a bottleneck in its memory. Using poor memory forces the Generator to learn a generalization of images and not memorize the images exactly.

```
# BUILD GENERATOR NETWORK
BadMemory = True

if BadMemory:
    seed = Input((10000,))
    x = Dense(2048, activation='elu')(seed)
    x = Reshape((8,8,32))(x)
    x = Conv2D(128, (3, 3), activation='elu', padding='same')(x)
    x = UpSampling2D((2, 2))(x)
    x = Conv2D(64, (3, 3), activation='elu', padding='same')(x)
    x = UpSampling2D((2, 2))(x)
    x = Conv2D(32, (3, 3), activation='elu', padding='same')(x)
    x = UpSampling2D((2, 2))(x)
    x = Conv2D(3, (3, 3), activation='linear', padding='same')(x)
    generated = Flatten()(x)
else:
    seed = Input((10000,))
```

```

generated = Dense(12288, activation='linear')(seed)

# COMPILE
generator = Model(seed, [generated, Reshape((10000,))(seed)])

# DISPLAY ARCHITECTURE
generator.summary()

```

Model: "functional_11"

Layer (type)	Output Shape	Param #	Connected to
input_layer_27 (InputLayer)	(None, 10000)	0	-
dense_2 (Dense)	(None, 2048)	20,482,048	input_layer_27[0][0]
reshape_2 (Reshape)	(None, 8, 8, 32)	0	dense_2[0][0]
conv2d (Conv2D)	(None, 8, 8, 128)	36,992	reshape_2[0][0]
up_sampling2d (UpSampling2D)	(None, 16, 16, 128)	0	conv2d[0][0]
conv2d_1 (Conv2D)	(None, 16, 16, 64)	73,792	up_sampling2d[0][0]
up_sampling2d_1 (UpSampling2D)	(None, 32, 32, 64)	0	conv2d_1[0][0]
conv2d_2 (Conv2D)	(None, 32, 32, 32)	18,464	up_sampling2d_1[0][0]
up_sampling2d_2 (UpSampling2D)	(None, 64, 64, 32)	0	conv2d_2[0][0]
conv2d_3 (Conv2D)	(None, 64, 64, 3)	867	up_sampling2d_2[0][0]
flatten_1 (Flatten)	(None, 12288)	0	conv2d_3[0][0]
reshape_3 (Reshape)	(None, 10000)	0	input_layer_27[0][0]

Total params: 20,612,163 (78.63 MB)

```

# BUILD GENERATIVE ADVERSARIAL NETWORK
model_discriminator.trainable = False

gan_input = Input(shape=(10000,))
x = generator(gan_input)
gan_output = model_discriminator(x)

gan = Model(gan_input, gan_output)

# Correct setting of conv2d weights
conv_layer = gan.get_layer('functional_11').get_layer('conv2d_1')
weights = conv_layer.get_weights()

# Set weights properly
new_kernel = np.zeros_like(weights[0])
new_bias = np.zeros_like(weights[1])
conv_layer.set_weights([new_kernel, new_bias])

# COMPILE GAN
gan.compile(optimizer=Adam(learning_rate=5.0), loss='mean_squared_error')

# DISPLAY ARCHITECTURE
gan.summary()

```

Model: "functional_22"

Layer (type)	Output Shape	Param #	Connected to
input_layer_38 (InputLayer)	(None, 10000)	0	-
functional_11 (Functional)	[(None, 12288), (None, 10000)]	20,612,163	input_layer_38[0][0]
functional (Functional)	(None, 12288)	122,892,290	functional_11[10][0], functional_11[10][1]

Total params: 143,504,453 (547.43 MB)

Trainable params: 20,612,163 (78.63 MB)

✓ Training GAN

```
train = np.zeros((10000,10000))
for i in range(10000): train[i,i] = 1
zeros = np.zeros((10000,12288))*500
```

```
Steps_per_epoch = 50
print("Training begins... Total steps per epoch: {}".format(Steps_per_epoch ))
```

➡ Training begins... Total steps per epoch: 50

```
# TRAIN NETWORKS
ep = 1; it = 9
if BadMemory: lr = 0.005
else: lr = 5.

for k in range(it):

    # BEGIN DISCRIMINATOR COACHES GENERATOR
    annealer = LearningRateScheduler(lambda x: lr)
    h = gan.fit(train, zeros, epochs = 10, batch_size=128, callbacks=[annealer], verbose=0)

    # DISPLAY GENERATOR LEARNING PROGRESS
    print('Epoch',(k+1),'/' +str(it)+' - loss =',h.history['loss'][-1] )
    plt.figure(figsize=(15,3))
    for j in range(5):
        xx = np.zeros((10000))
        xx[np.random.randint(10000)] = 1
        plt.subplot(1,5,j+1)
        img = generator.predict(xx.reshape((-1, 10000)))[0]





        # Ensure the image is reshaped to (64, 64, 3)
        img = img.reshape((64, 64, 3))

        # Replace NaNs or Infs with 0
        img = np.nan_to_num(img, nan=0.0, posinf=255, neginf=0)





        # Clip image values to be within the valid range for uint8 (0-255)
        img = np.clip(img, 0, 255).astype('uint8')
        img_pil = Image.fromarray(img)
        plt.axis('off')
        plt.imshow(img)
    plt.show()

    # ADJUST LEARNING RATES
    if BadMemory:
        ep *= 2
        if ep>=32: lr = 0.001
        if ep>256: ep = 256
    else:
        if h.history['loss'][-1] < 25: lr = 1.
        if h.history['loss'][-1] < 1.5: lr = 0.5
```






Epoch 1 / 9 - loss = 0.0
1/1 0s 17ms/step
1/1 0s 18ms/step
1/1 0s 21ms/step
1/1 0s 17ms/step
1/1 0s 18ms/step



Epoch 2 / 9 - loss = 0.0
1/1 0s 20ms/step
1/1 0s 19ms/step
1/1 0s 18ms/step
1/1 0s 17ms/step
1/1 0s 16ms/step



Epoch 3 / 9 - loss = 0.0
1/1 0s 21ms/step
1/1 0s 17ms/step
1/1 0s 16ms/step
1/1 0s 17ms/step
1/1 0s 16ms/step



```
-----  
KeyboardInterrupt                                Traceback (most recent call last)  
/tmp/ipykernel_31/2868487983.py in <cell line: 0>()  
      8     # BEGIN DISCRIMINATOR COACHES GENERATOR  
      9     annealer = LearningRateScheduler(lambda x: lr)  
     10     # BEGIN DISCRIMINATOR COACHES GENERATOR
```